

# Introduction to Amazon Athena

Interactive, Serverless, Pay-per-use, Query Service

Presented by:

**Raul Frias**

Solutions Architect Manager

**Kevin Epstein**

Head of Cloud Acceleration

# What to Expect from the Session

- Overview of Amazon Athena
- Key Features
- Customer Examples
- Troubleshooting Query errors
- Q&A

# Challenges Customers Faced


- Significant amount of work required to analyze data in Amazon S3
- Users often only have access to aggregated data sets
- Managing a Hadoop cluster or data warehouse requires expertise

# Introducing Amazon Athena

Amazon Athena is an **interactive query service** that makes it easy to analyze data directly from Amazon S3 using Standard SQL

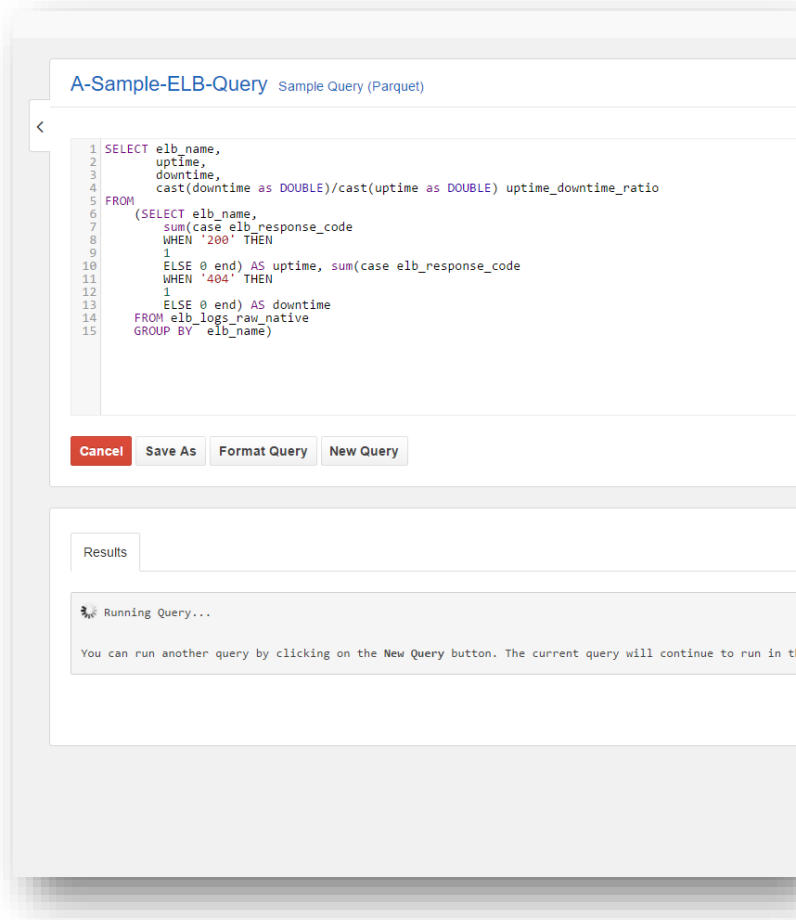
# Athena is Serverless

- No Infrastructure or administration
- Zero Spin up time
- Transparent upgrades



# Amazon Athena is Easy To Use

- Log into the Console
- Create a table
  - Type in a Hive DDL Statement
  - Use the console Add Table wizard
- Start querying



# Amazon Athena is Highly Available

- You connect to a service endpoint or log into the console
- Athena uses warm compute pools across multiple Availability Zones
- Your data is in Amazon S3, which is also highly available and designed for 99.999999999% durability

# Query Data Directly from Amazon S3

- No loading of data
- Query data in its raw format
  - Text, CSV, JSON, weblogs, AWS service logs
  - Convert to an optimized form like ORC or Parquet for the best performance and lowest cost
- No ETL required
- Stream data directly from Amazon S3
- Take advantage of Amazon S3 durability and availability



# Use ANSI SQL

- Start writing ANSI SQL
- Support for complex joins, nested queries & window functions
- Support for complex data types (arrays, structs)
- Support for partitioning of data by any key
  - (date, time, custom keys)
  - e.g., Year, Month, Day, Hour or Customer Key, Date

```
1 WITH q21_tmp1_cached AS
2 (SELECT l_orderkey,
3      count(DISTINCT l_suppkey) AS count_suppkey,
4      max(l_suppkey) AS max_suppkey
5 FROM lineitem_parq
6 WHERE l_orderkey IS NOT NULL
7 GROUP BY l_orderkey),
8 q21_tmp2_cached AS
9 (SELECT l_orderkey,
10      count(DISTINCT l_suppkey) count_suppkey,
11      max(l_suppkey) AS max_suppkey
12 FROM lineitem_parq
13 WHERE l_receiptdate > l_commitdate
14 AND l_orderkey IS NOT NULL
15 GROUP BY l_orderkey)
16 SELECT s_name,
17      count(1) AS numwait
18 FROM
19 (SELECT s_name
20 FROM
21      (SELECT s_name,
22           t2.l_orderkey,
23           l_suppkey,
24           count_suppkey,
25           max_suppkey
26 FROM q21_tmp2_cached t2
27 RIGHT OUTER JOIN
28      (SELECT s_name,
29           l_orderkey,
30           l_suppkey
31 FROM
32      (SELECT s_name,
33           t1.l_orderkey,
34           l_suppkey,
35           count_suppkey,
36           max_suppkey
37 FROM q21_tmp1_cached t1
38 JOIN
39      (SELECT s_name,
40           l_orderkey,
41           l_suppkey
42 FROM orders_parq o
43 JOIN
44      (SELECT s_name,
45           l_orderkey,
46           l_suppkey
47 FROM nation_parq n
48 JOIN supplier s ON s.s_nationkey = n.n_nationkey
49 AND n.n_name = 'SAUDI ARABIA'
50 JOIN lineitem_parq l ON s.s_suppkey = l.l_suppkey
51 WHERE l.l_receiptdate > l.l_commitdate
52 AND l.l_orderkey IS NOT NULL) t1 ON o.o_orderkey = t1.l_orderkey
53 AND o.o_orderstatus = 'F') t2 ON t2.l_orderkey = t1.l_orderkey) a
54 WHERE (count_suppkey > 1)
55 OR ((count_suppkey=1)
56 AND (l_suppkey <= max_suppkey))) t3 ON t3.l_orderkey = t2.l_orderkey) b
57 WHERE (count_suppkey IS NULL)
58 OR ((count_suppkey=1)
59 AND (l_suppkey = max_suppkey))) c
60 GROUP BY s_name
61 ORDER BY numwait DESC,
62      s_name LIMIT 100;
```

# Familiar Technologies Under the Covers



## Used for SQL Queries

In-memory distributed query engine  
ANSI-SQL compatible with extensions



## Used for DDL functionality


Complex data types  
Multitude of formats  
Supports data partitioning

# Amazon Athena Supports Multiple Data Formats

- Text files, e.g., CSV, raw logs
- Apache Web Logs, TSV files
- JSON (simple, nested)
- Compressed files
- Columnar formats such as Apache Parquet & Apache ORC
- AVRO support – coming soon

# Amazon Athena is Fast

- Tuned for performance
- Automatically parallelizes queries
- Results are streamed to console
- Results also stored in S3
- Improve Query performance
  - Compress your data
  - Use columnar formats



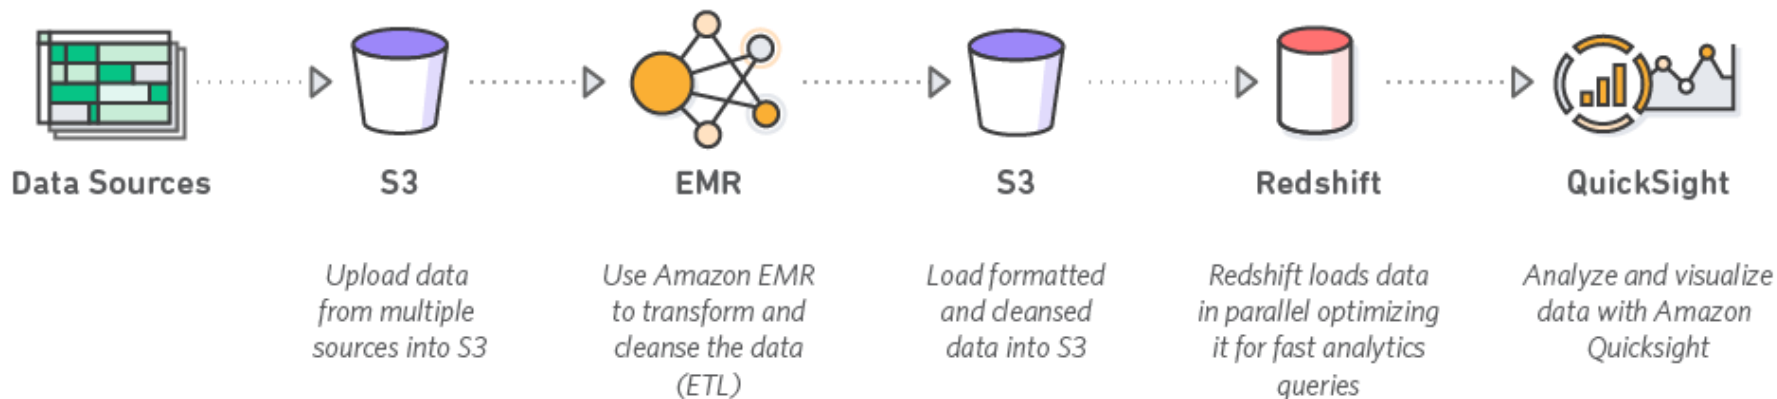
# Amazon Athena is Cost Effective

- Pay per query
- \$5 per TB scanned from S3
- DDL Queries and failed queries are free
- Save by using compression, columnar formats, partitions

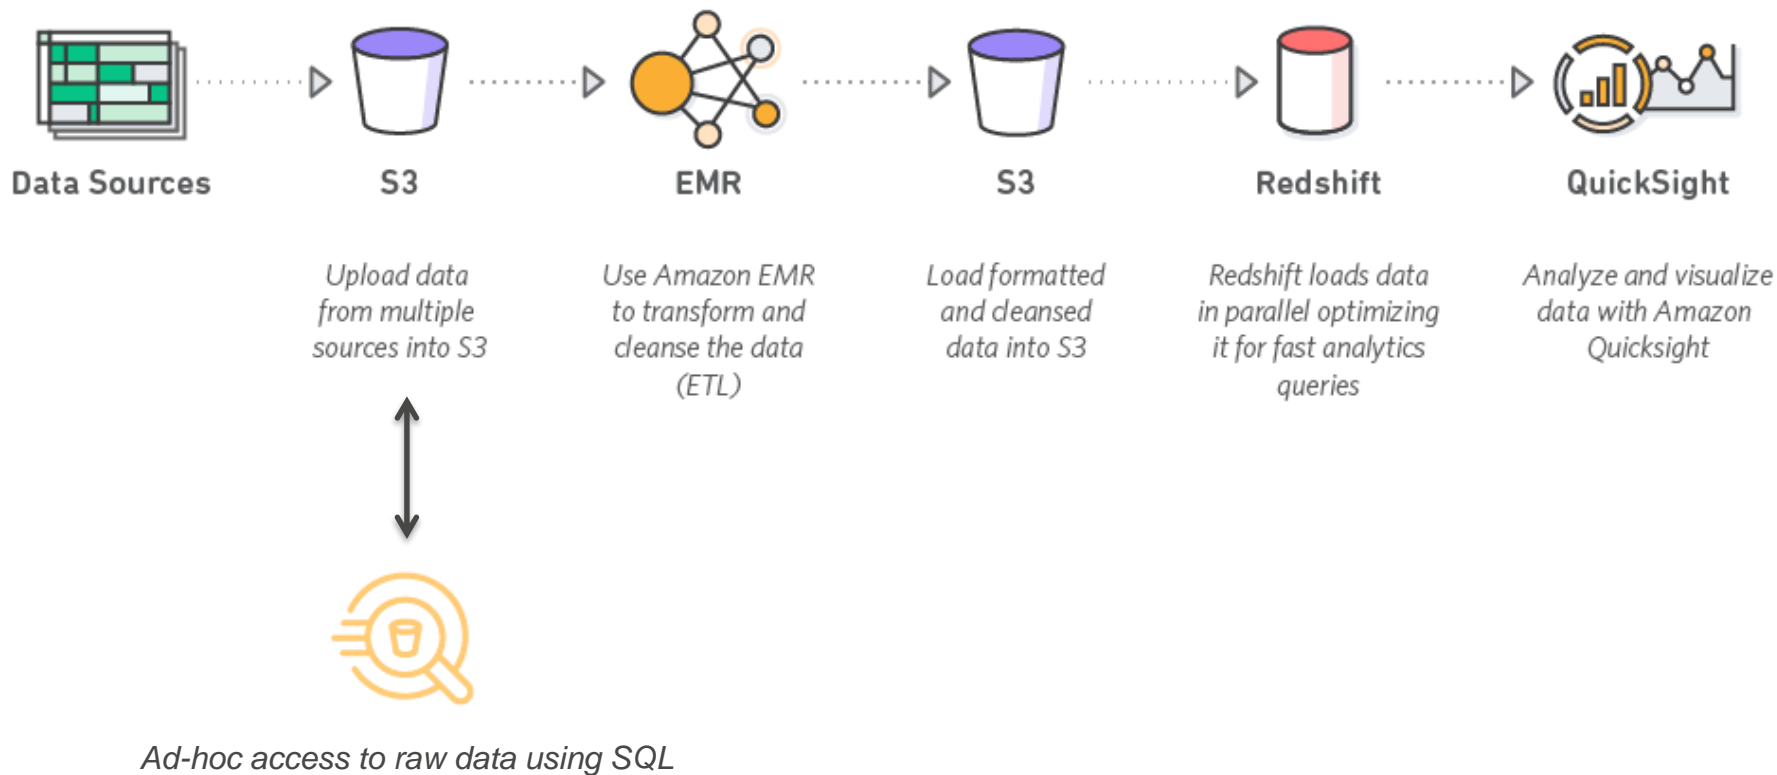
# Who is Athena for ?

- Any one looking to process data stored in Amazon S3
  - Data coming IOT Devices, Apache Logs, Omniture logs, CF logs, Application Logs
- Anyone who knows SQL
  - Both developers or Analysts
- Ad-hoc exploration of data and data discovery
- Customers looking to build a data lake on Amazon S3

# A Sample Pipeline

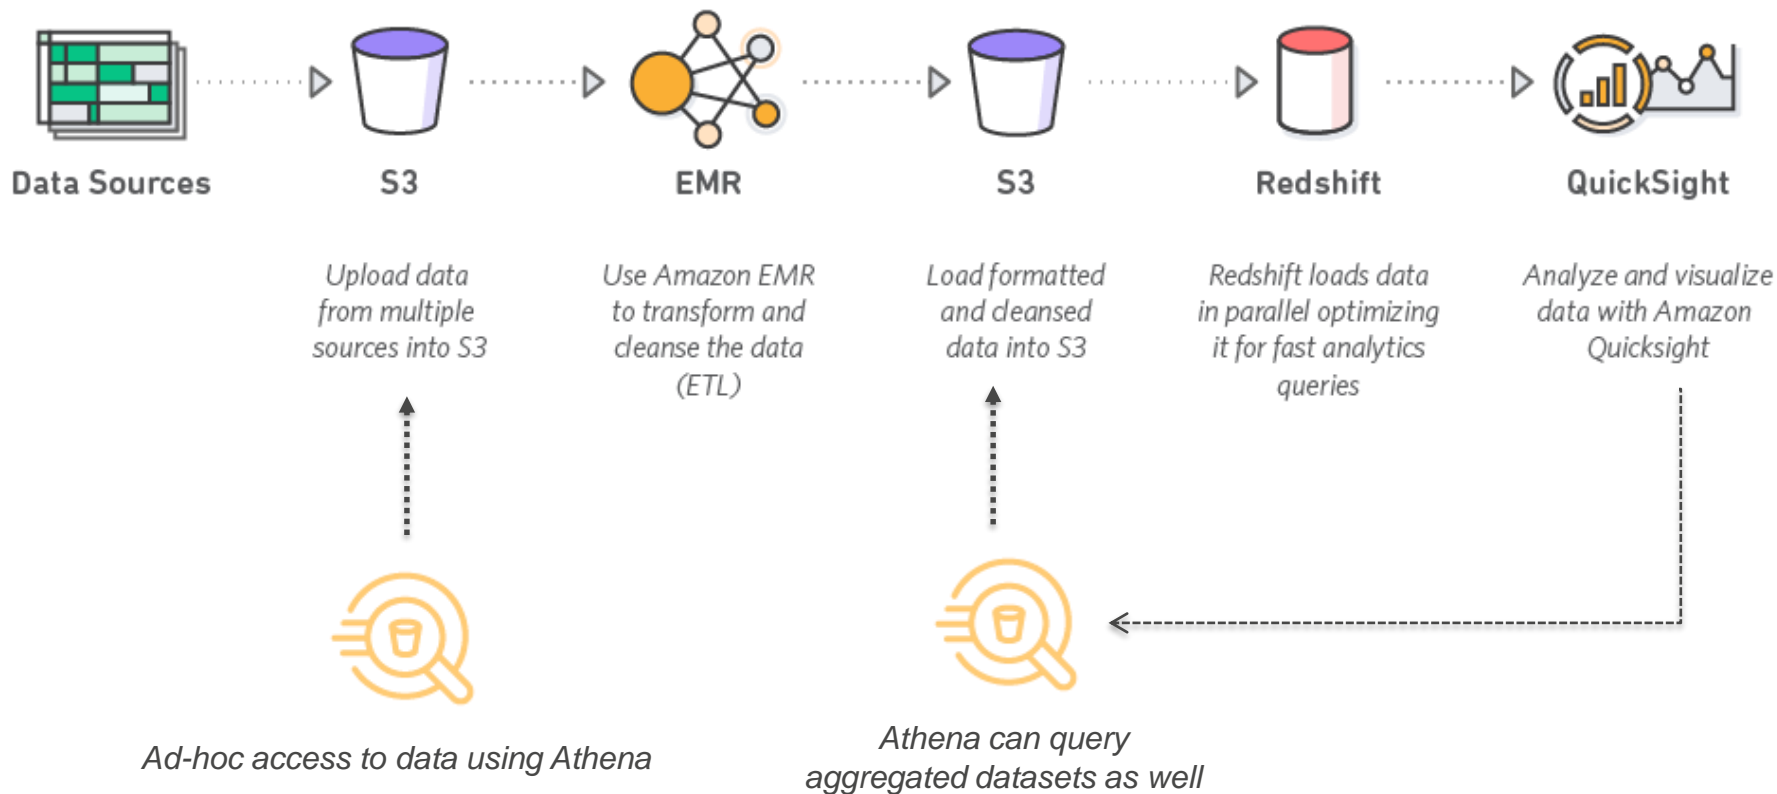


# A Sample Pipeline





# A Sample Pipeline



# Re-visiting Challenges

~~Significant amount of work required to analyze data in Amazon S3~~

No ETL required. No loading of data. Query data where it lives

~~Users often only have access to aggregated data sets~~

Query data at whatever granularity you want

~~Managing a Hadoop cluster or data warehouse requires expertise~~

No infrastructure to manage

# Accessing Amazon Athena

# Use the JDBC Driver

Select Connection Profile

Default group

Athena

Driver: Athena (com.amazonaws.athena.jdbc.AthenaDriver)

URL: jdbc:awsathena://athena.us-east-1.amazonaws.com:443

Username:

Password:  Show password

Autocommit: ☐ Fetch size:  Timeout:  s Extended Properties

☐ Prompt for username ☐ Confirm updates ☐ Read only ☒ Remember DbExplorer Schema

☒ Save password ☐ Confirm DML without WHERE ☐ Store completion cache locally

☒ Separate connection per tab ☐ Rollback before disconnect ☐ Remove comments

☐ Ignore DROP errors ☐ Empty string is NULL ☐ Hide warnings

☐ Trim CHAR data ☒ Include NULL columns in INSERTS ☐ Check for uncommitted changes

Info Background: ☐ X ... (None) Alternate Delimiter:

Workspace:  ...

Main window icon:  ...

Macros:  ...

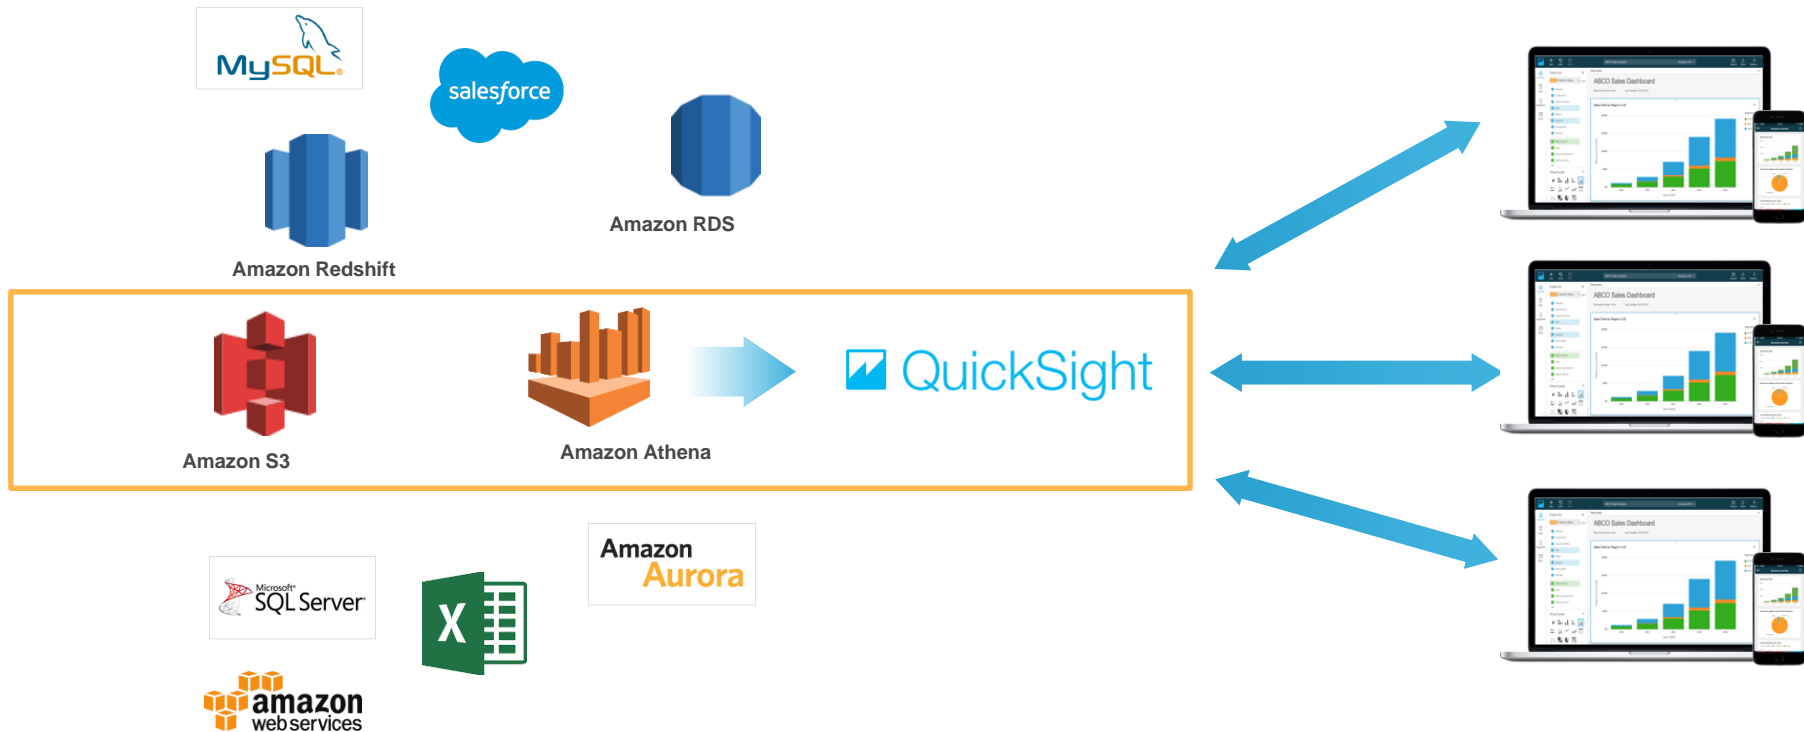
Tags:

Connect scripts Schema/Catalog Filter Variables Test

Columns SQL source Data Indexes References Refere								
COLUMN_NAME								
COLUMN_NAME	DATA_TYPE	PK	NULLABLE	DEFAULT	AUTOINCREMENT	COMPUTED	REMARKS	JDBC Type
date	date	NO	NO		NO	NO		DATE
time	string	NO	NO		NO	NO		LONGNVARCHAR
location	string	NO	NO		NO	NO		LONGNVARCHAR
bytes	int	NO	NO		NO	NO		INTEGER
requestip	string	NO	NO		NO	NO		LONGNVARCHAR
method	string	NO	NO		NO	NO		LONGNVARCHAR
host	string	NO	NO		NO	NO		LONGNVARCHAR
uri	string	NO	NO		NO	NO		LONGNVARCHAR
status	int	NO	NO		NO	NO		INTEGER
referrer	string	NO	NO		NO	NO		LONGNVARCHAR
os	string	NO	NO		NO	NO		LONGNVARCHAR

# Using Amazon Athena with Amazon QuickSight


QuickSight allows you to connect to data from a wide variety of AWS, third-party, and on-premises sources including Amazon Athena




Data sets

79.8MB of SPICE used of 141GB in N. Virginia


Create a Data Set  
FROM NEW DATA SOURCES




Upload a file  
(.csv, .tsv, .clf, .elf, .xlsx)




Salesforce  
Connect to Salesforce




S3




RDS




Redshift  
Auto-discovered




Redshift  
Manual connect




Athena




MySQL




PostgreSQL



SQL Server



Aurora



MariaDB


FROM EXISTING DATA SOURCES



People Overview  
Updated 2 days ago



Web and Social Media A...  
Updated 2 days ago



Sales Pipeline  
Updated 2 days ago

Data sets

## New Athena data source



## Data source name

Create data source

79.8MB of SPICE used of 141GB in N. Virginia

## Create a Data Set

FROM NEW DATA SOURCES



## Upload a file

(.csv, .tsv, .cif, .elf, .xlsx)



## Salesforce

Connect to Salesforce



## S3



## RDS



## Redshift

Auto-discovered



## Redshift

Manual connect



## Athena



## MySQL



## PostgreSQL



## SQL Server



## Aurora



## MariaDB

FROM EXISTING DATA SOURCES



## People Overview

Updated 2 days ago



## Web and Social Media A...

Updated 2 days ago



## Sales Pipeline

Updated 2 days ago

Data sets

Create a Data Set  
FROM NEW DATA SOURCES



Upload a file  
(.csv, .tsv, .clf, .elf, .xlsx)



RDS

Athena

SQL Server

Sample Data

Updated a month ago

Choose your table



My Athena Connection

Database: contain sets of tables.

Flight data



Tables: contain the data you can visualize.

- ☐ airport\_id
- ☒ all\_flights
- ☐ cancellation\_id
- ☐ carrier\_id
- ☐ delay\_id

Edit/Preview data

Select

test1111

Updated a month ago

47lining data

Updated a month ago



Data sets



RDS



Athena



SQL Server



Aurora



MariaDB



Sample Data

Updated a month ago

test1111

Updated a month ago



47lining data

Updated a month ago



Sample Data

Updated a month ago

SN-Cluster

Updated a month ago



47lining data

Updated a month ago

## Finish data set creation



Database: Flight data  
Table: all\_flights  
Data source: My Athena Connection

- ☐ Import to SPICE
- ☒ Directly query your data

✓ 248.3GB available

SPICE

Edit/Preview data

Visualize

21.7GB of SPICE used of 270GB in N. Virginia

## FROM EXISTING DATA SOURCES



## Fields list

all\_flights

- # day\_of\_week
- # dep\_time
- # dest\_airport\_id
- # dest\_city\_market\_id
- # dest\_city\_name
- # dest\_state\_abr
- # dest\_state\_nm
- # dest\_wac
- # fl\_date

- # cancelled
- # carrier\_delay
- # dep\_del15
- # dep\_delay
- # dep\_delay\_group


## Visual types




## Field wells

### Total flights by month


Viewing 2...



### Most popular destinations



### Flights by day of week



# JDBC also Provides Programmatic Access

```
/* Setup the driver */
```

```
Properties info = new Properties();  
info.put("user", "AWSAccessKey");  
info.put("password", "AWSSecretAccessKey");  
info.put("s3_staging_dir", "s3://S3 Bucket Location/");
```

```
Class.forName("com.amazonaws.athena.jdbc.AthenaDriver");
```

```
Connection connection = DriverManager.getConnection("jdbc:awsathena://athena.us-east-1.amazonaws.com:443/", info);
```

# Creating a Table and Executing a Query

*/\* Create a table \*/*

```
Statement statement = connection.createStatement();
```

```
ResultSet queryResults = statement.executeQuery("CREATE EXTERNAL TABLE tableName ( Col1  
String ) LOCATION 's3://bucket/tableLocation'");
```

*/\* Execute a Query \*/*

```
Statement statement = connection.createStatement();
```

```
ResultSet queryResults = statement.executeQuery("SELECT * FROM cloudfront_logs");
```

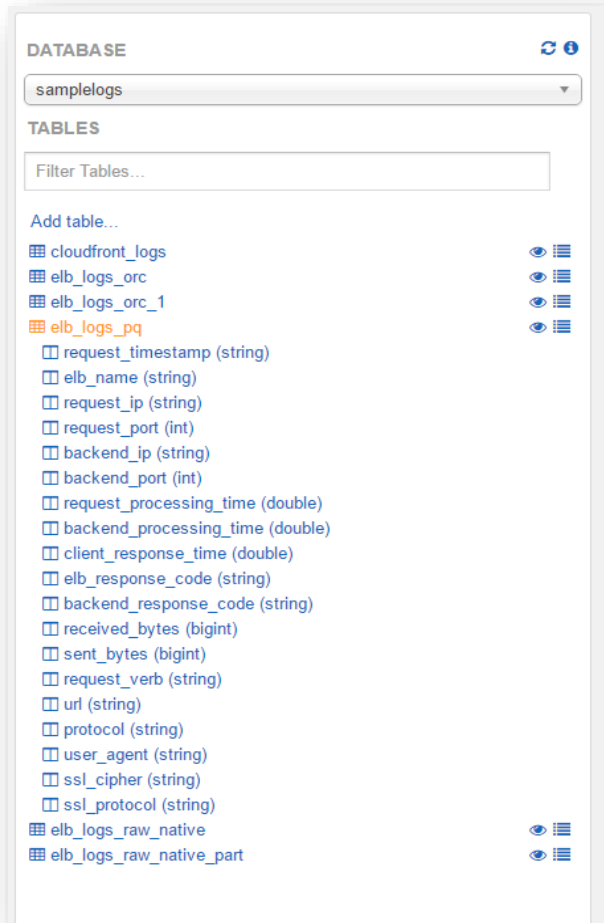
# Creating Tables and Querying Data

# Creating Tables - Concepts

- Create Table Statements (or DDL) are written in Hive
  - High degree of flexibility
  - Schema on Read
  - Hive is SQL like but allows other concepts such “external tables” and partitioning of data
  - Data formats supported – JSON, TXT, CSV, TSV, Parquet and ORC (via Serdes)
  - Data is stored in Amazon S3
  - Metadata is stored in a metadata store

# Athena's Internal Metadata Store

- Stores Metadata
  - Table definition, column names, partitions
- Highly available and durable
- Requires no management
- Access via DDL statements
- Similar to a Hive Metastore



# Running Queries is Simple

```
1 SELECT o_year,  
2        sum(case  
3          WHEN nation = 'PERU' THEN  
4            volume  
5          ELSE 0 end) / sum(volume) AS mkt_share  
6 FROM  
7   (SELECT substr(o_orderdate,  
8     1,  
9     4) AS o_year,  
10    l_extendedprice * (1 - l_discount) AS volume,  
11    n2.n_name AS nation  
12 FROM lineitem_parq, orders_parq, part_parq, customer_parq, supplier_parq, nation_parq n1, nation_parq n2, region_parq  
13 WHERE p_partkey = l_partkey  
14        AND s_suppkey = l_suppkey  
15        AND l_orderkey = o_orderkey  
16        AND o_custkey = c_custkey  
17        AND c_nationkey = n1.n_nationkey  
18        AND n1.n_regionkey = r_regionkey  
19        AND r_name = 'AMERICA'  
20        AND s_nationkey = n2.n_nationkey  
21        AND o_orderdate  
22        BETWEEN '1995-01-01'  
23        AND '1996-01-01')
```

Run Query

Save As

Format Query

New Query

(Run time: 17.34 seconds, Data scanned: 12.97GB)

Run time  
and data  
scanned

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

	o_year	mkt_share
1	1995	0.04048235539866028
2	1996	0.03984664161294089



# Apache Parquet and Apache ORC – Columnar Formats

## PARQUET

- Columnar format
- Schema segregated into footer
- Column major format
- All data is pushed to the leaf
- Integrated compression and indexes
- Support for predicate pushdown

## ORC

- Apache Top level project
- Schema segregated into footer
- Column major with stripes
- Integrated compression, indexes, and stats
- Support for Predicate Pushdown

# Converting to ORC and PARQUET

- You can use Hive CTAS to convert data
  - CREATE TABLE new\_key\_value\_store
  - STORED AS PARQUET
  - AS
  - SELECT col\_1, col2, col3 FROM noncolumartable
  - SORT BY new\_key, key\_value\_pair;
- You can also use Spark to convert the file into PARQUET / ORC
- 20 lines of Pyspark code, running on EMR
  - Converts 1TB of text data into 130 GB of Parquet with snappy conversion
  - Total cost \$5

<https://github.com/awslabs/aws-big-data-blog/tree/master/aws-blog-spark-parquet-conversion>

# Pay By the Query - \$5/TB Scanned

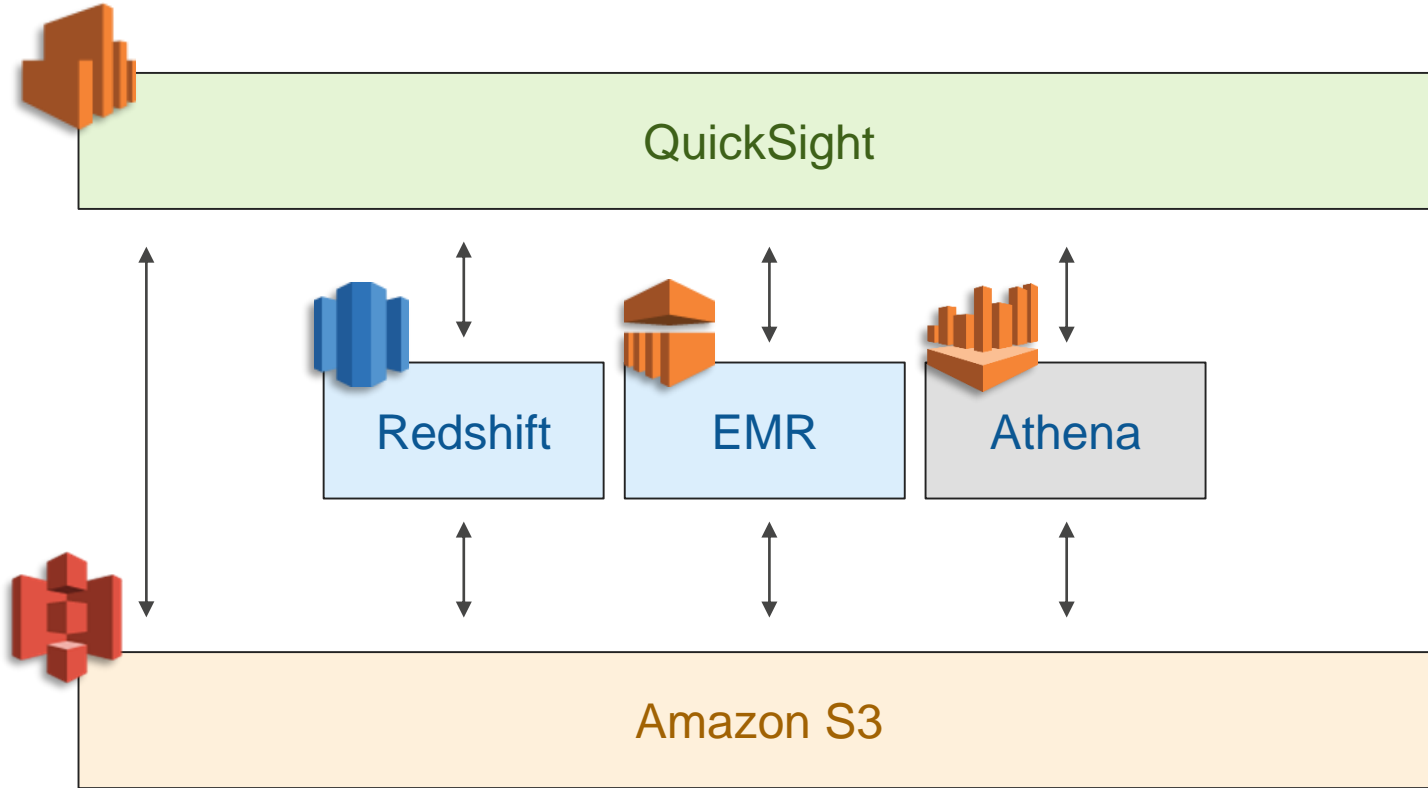
- Pay by the amount of data scanned per query
- Ways to save costs
  - Compress
  - Convert to Columnar format
  - Use partitioning
- Free: DDL Queries, Failed Queries

```
SELECT elb_name,
       uptime,
       downtime,
       cast(downtime as DOUBLE)/cast(uptime as DOUBLE) uptime_downtime_ratio
FROM
  (SELECT elb_name,
         sum(case elb_response_code
                WHEN '200' THEN
                1
                ELSE 0 end) AS uptime, sum(case elb_response_code
                WHEN '404' THEN
                1
                ELSE 0 end) AS downtime
    FROM elb_logs_raw_native
   GROUP BY elb_name)
```

Dataset	Size on Amazon S3	Query Run time	Data Scanned	Cost
Logs stored as Text files	1 TB	237 seconds	1.15TB	\$5.75
Logs stored in Apache Parquet format*	130 GB	5.13 seconds	2.69 GB	\$0.013
Savings	87% less with Parquet	34x faster	99% less data scanned	99.7% cheaper

# Use Cases

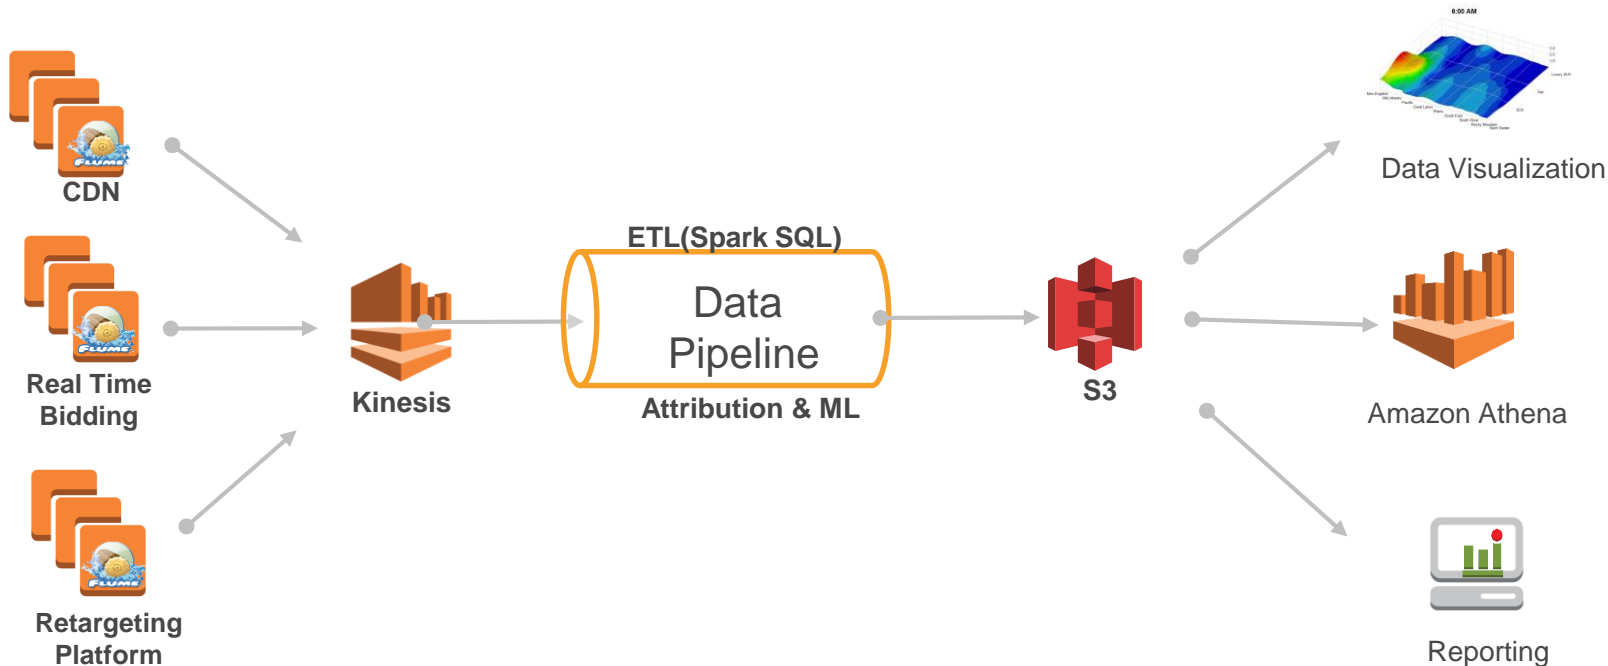
# Athena Complements Amazon Redshift & Amazon EMR



# Customers Using Athena



# DataXu – 180TB of Log Data per Day





**Rob Harrop**

@robertharrop



Following

Up and running with AWS Athena already,  
querying production performance data from logs  
in S3.

RETWEET

1

LIKES

4



10:24 AM - 30 Nov 2016



1



1



4





# Tips and Tricks

# Created a Table, but do not see data

- Verify that the input LOCATION is correct on S3
  - Buckets should be specified as `s3://name/` or `s3://name/subfolder/`
  - `s3://us-east-1.amazonaws.com/bucket/path` will throw an error `s3://bucket/path/`
- Did the table have partitions ?
  - MSCK Repair Table for
    - `s3://mybucket/athena/inputdata/year=2016/data.csv`
    - `s3://mybucket/athena/inputdata/year=2015/data.csv`
    - `s3://mybucket/athena/inputdata/year=2014/data.csv`
  - Alter Table Add Partition for
    - `s3://mybucket/athena/inputdata/2016/data.csv`
    - `s3://mybucket/athena/inputdata/2015/data.csv`
    - `s3://mybucket/athena/inputdata/2014/data.csv`
  - ALTER TABLE Employee ADD
    - `PARTITION (year=2016) LOCATION s3://mybucket/athena/inputdata/2016/`
    - `PARTITION (year=2015) LOCATION s3://mybucket/athena/inputdata/2015/`
    - `PARTITION (year=2014) LOCATION s3://mybucket/athena/inputdata/2014/`

# How did you define your partitions

```
CREATE EXTERNAL TABLE Employee (  
    Id INT,  
    Name STRING,  
    Address STRING  
) PARTITIONED BY (year INT)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY ','  
LOCATION  
's3://mybucket/athena/inputdata/';
```

```
CREATE EXTERNAL TABLE Employee (  
    Id INT,  
    Name STRING,  
    Address STRING,  
    INT Year  
) PARTITIONED BY (year INT)  
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY ','  
LOCATION  
's3://mybucket/athena/inputdata/';
```

# Reading JSON Data

- Make sure you are using the right Serde
  - Native JSON Serde `org.apache.hive.hcatalog.data.JsonSerDe`
  - OpenX SerDe (`org.openx.data.jsonserde.JsonSerDe`)
- Make sure JSON record is a single line
- Generate your data in case-insensitive columns
- Provide an option to ignore malformed records

```
CREATE EXTERNAL TABLE json (  
    a string,  
    b int  
)  
ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'  
WITH SERDEPROPERTIES ( 'ignore.malformed.json' = 'true')  
LOCATION 's3://bucket/path/';
```

# Access Denied Issues

Check the IAM Policy

Refer to the Getting Started Documentation

Check the bucket ACL

Both the read bucket and write bucket

# Table names

- Use backticks if table names begin with an underscore.

For example:

```
CREATE TABLE myUnderScoreTable (  
  `_id` string,  
  `_index` string,
```

- Table name can only have underscores as special characters

# Thank you!